

Socket Programming


Rathindra Nath Dutta

Senior Research Fellow
Advanced Computing & Microelectronics Unit
Indian Statistical Institute, Kolkata



October 21, 2022

https://www.isical.ac.in/~rathin_r/uploads/CN/2022/Socket.html



WEB PAGE

What is a Socket?

- “A network **socket** is a software structure within a network node of a computer network that serves as an endpoint for sending and receiving data across the network ...Sockets are created only during the lifetime of a process of an application running in the node” – WIKI¹
- “A **socket** is a communications connection point (endpoint) that you can name and address in a network. Socket programming shows how to use socket APIs to establish communication links between remote and local processes” – IBM²
- “A **socket** is one endpoint of a two-way communication link between two programs running on the network” – ORACLE³

¹https://en.wikipedia.org/wiki/Network_socket

²<https://www.ibm.com/docs/en/i/7.5?topic=communications-socket-programming>

³<https://docs.oracle.com/javase/tutorial/networking/sockets/definition.html>

How are Sockets uniquely identified?

- A *Socket* is identified to other hosts by its **Socket Address**

How are Sockets uniquely identified?

- A *Socket* is identified to other hosts by its **Socket Address**
- A *Socket Address* consists of a **Transport Protocol** (TCP, UDP, etc.), an **IP Address** (IPv4 or IPv6) and a **Port Number**

How are Sockets uniquely identified?

- A *Socket* is identified to other hosts by its **Socket Address**
- A *Socket Address* consists of a **Transport Protocol** (TCP, UDP, etc.), an **IP Address** (IPv4 or IPv6) and a **Port Number**
- A *Port Number* is a (unique) 16-bit unsigned integer assigned to one endpoint

How are Sockets uniquely identified?

- A *Socket* is identified to other hosts by its **Socket Address**
- A *Socket Address* consists of a **Transport Protocol** (TCP, UDP, etc.), an **IP Address** (IPv4 or IPv6) and a **Port Number**
- A *Port Number* is a (unique) 16-bit unsigned integer assigned to one endpoint
- ports 0 through 1023 are well-known ports (aka system ports)
- ports 1024 through 49151 are registered ports⁴
- ports from 49152 through 65535 are dynamic or private ports; commonly known as ephemeral ports⁵

⁴IANA maintains the official list https://en.wikipedia.org/wiki/List_of_TCP_and_UDP_port_numbers

⁵ephemeral (adj.): lasting for a very short time or having a very short life cycle

Why another address?

TCP/IP Layer

Application

Transport

Network

Data-Link
Physical

Why another address?

TCP/IP Layer	Common Protocols
Application	TELNET, HTTP, DHCP, PING, FTP, ...
Transport	TCP, UDP, ...
Network	IP, ARP, TCMP, ...
Data-Link Physical	Ethernet, WiFi, ...

Why another address?

TCP/IP Layer	Common Protocols	Data Packet
Application	TELNET, HTTP, DHCP, PING, FTP, ...	Message
Transport	TCP, UDP, ...	Segment/Datagram
Network	IP, ARP, TCMP, ...	Datagram
Data-Link Physical	Ethernet, WiFi, ...	Frame Bits

Why another address?

TCP/IP Layer	Common Protocols	Data Packet	Address
Application	TELNET, HTTP, DHCP, PING, FTP, ...	Message	Application Specific
Transport	TCP, UDP, ...	Segment/Datagram	Port
Network	IP, ARP, TCMP, ...	Datagram	Logical (IP)
Data-Link Physical	Ethernet, WiFi, ...	Frame Bits	Physical(MAC)

Why another address?

TCP/IP Layer	Common Protocols	Data Packet	Address	Objective
Application	TELNET, HTTP, DHCP, PING, FTP, ...	Message	Application Specific	
Transport	TCP, UDP, ...	Segment/Datagram	Port	
Network	IP, ARP, TCMP, ...	Datagram	Logical (IP)	
Data-Link Physical	Ethernet, WiFi, ...	Frame Bits	Physical(MAC)	identification

Why another address?

TCP/IP Layer	Common Protocols	Data Packet	Address	Objective
Application	TELNET, HTTP, DHCP, PING, FTP, ...	Message	Application Specific	
Transport	TCP, UDP, ...	Segment/Datagram	Port	
Network	IP, ARP, TCMP, ...	Datagram	Logical (IP)	logical organization ¹
Data-Link Physical	Ethernet, WiFi, ...	Frame Bits	Physical(MAC)	identification

¹better management, efficient routing

Why another address?

TCP/IP Layer	Common Protocols	Data Packet	Address	Objective
Application	TELNET, HTTP, DHCP, PING, FTP, ...	Message	Application Specific	
Transport	TCP, UDP, ...	Segment/Datagram	Port	host-to-host delivery ²
Network	IP, ARP, TCMP, ...	Datagram	Logical (IP)	logical organization ¹
Data-Link Physical	Ethernet, WiFi, ...	Frame Bits	Physical(MAC)	identification

²process-to-process, a node can run multiple processes each talking via different protocol

¹better management, efficient routing

Why another address?

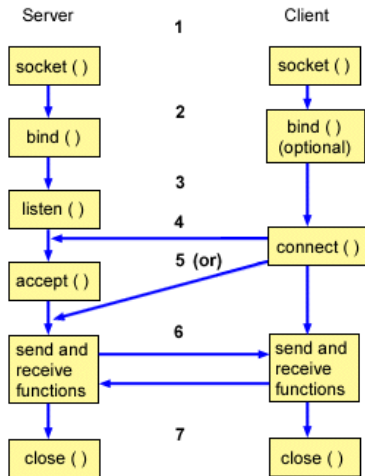
TCP/IP Layer	Common Protocols	Data Packet	Address	Objective
Application	TELNET, HTTP, DHCP, PING, FTP, ...	Message	Application Specific	end-to-end delivery ³
Transport	TCP, UDP, ...	Segment/Datagram	Port	host-to-host delivery ²
Network	IP, ARP, TCMP, ...	Datagram	Logical (IP)	logical organization ¹
Data-Link Physical	Ethernet, WiFi, ...	Frame Bits	Physical(MAC)	identification

³session control, a process(browser) may have multiple active session(tabs) to same client(google search), uses application specific URIs

²process-to-process, a node can run multiple processes each talking via different protocol

¹better management, efficient routing

Socket Programming Using C



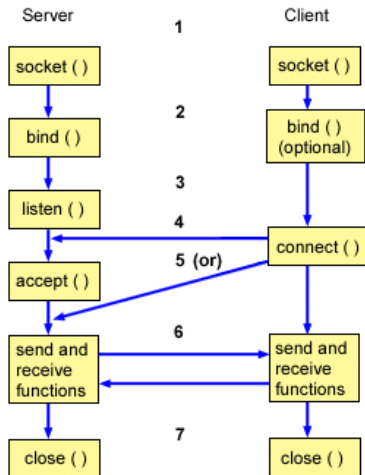
`socket()` creates and returns a socket descriptor representing an end-point for communications

Servers must bind a unique name to a socket descriptor using `bind()` to make it accessible from the network

`listen()` call shows willingness to accept client connection requests
NB: a socket cannot actively initiate any connection requests after a `listen()` call

image src: <https://www.ibm.com/docs/en/i/7>.

Socket Programming Using C




The client invokes `connect()` on a stream socket to establish a connection to the server

The server uses `accept()` to accept a client connection request
NB: The server must issue `bind()` and `listen()` calls successfully before `accept()`

image src: <https://www.ibm.com/docs/en/i/7>.

5?topic=programming-how-sockets-work

Socket Programming Using C



When a connection is established between stream sockets (between client and server), we can use any of the data transfer methods of socket APIs such as `send()`, `recv()`, `read()`, `write()`, ...

Finally, when a server or client wants to stop operations, it issues a `close()` call to release any system resources acquired by the socket

image src: <https://www.ibm.com/docs/en/i/7>.

The `socket()` API

```
#include <sys/socket.h>
```

```
int socket(int domain, int type, int protocol);
```

Return value: On success, a file descriptor for the new socket is returned, lowest-numbered file descriptor not currently open for the process. On error, -1 is returned

Parameters:

domain: specifies a communication domain; selects the protocol family which will be used for communication

common domain values: `AF_UNIX` (Local communication), `AF_INET` (IPv4 Internet protocols), `AF_INET6` (IPv6 Internet protocols)

type: specifies the communication semantics commonly used types are:

`SOCK_STREAM` (sequenced, reliable, two-way, connection-oriented byte streams, TCP),

`SOCK_DGRAM` (connectionless, unreliable messages of a fixed maximum length, UDP)

protocol: specifies a particular protocol to be used with the socket usually a 0 is specified to denote the default protocol for the corresponding socket type

The bind() API

```
#include <sys/socket.h>
int bind(int sockfd, const struct sockaddr *addr,
         socklen_t addrlen);
```

assigns the address specified by `addr` to the socket referred to by the file descriptor `sockfd`

Return value: On success, zero is returned. On error, -1 is returned

Parameters:

sockfd: a socket file descriptor created with `socket()`

addr: a pointer to an address structure, actual structure depends on the socket address family

addrlen: specifies the size, in bytes, of the address structure pointed to by `addr`

Address Structure for AF_INET

```
#include <sys/socket.h>
#include <netinet/in.h>
struct sockaddr_in {
    sa_family_t    sin_family; /* address family: AF_INET */
    in_port_t      sin_port;   /* port in network byte order */
    struct in_addr sin_addr;   /* internet address */
};

/* Internet address */
struct in_addr {
    uint32_t        s_addr;    /* address in network byte order */
};
```

¹<https://man7.org/linux/man-pages/man7/ip.7.html>

Binding a Socket to an Address

```
struct sockaddr_in sock_addr;  
bzero((char *)&sock_addr, sizeof(sock_addr)); //clear  
sockfd = socket(AF_INET, SOCK_STREAM, 0);  
int portno = 54321;  
  
sock_addr.sin_family = AF_INET;  
sock_addr.sin_port = htons(portno);  
sock_addr.sin_addr.s_addr = INADDR_ANY;  
//sock_addr.sin_addr.s_addr = INADDR_LOOPBACK  
//sock_addr.sin_addr.s_addr = inet_addr("127.0.0.1");  
bind(sockfd, (struct sockaddr*)&sock_addr, sizeof(sock_addr))
```

¹htons(): converts an unsigned short integer from host byte order to network byte order

²special addresses: INADDR_LOOPBACK (127.0.0.1) always refers to the *localhost* via the loopback device;
INADDR_ANY (0.0.0.0) means any address for binding; INADDR_BROADCAST (255.255.255.255) means any host

³inet_addr(): converts a IPv4 host address string written in dotted decimal notation, into binary data
in network byte order; require #include <arpa/inet.h>

The listen() API

```
#include <sys/socket.h>
```

```
int listen(int sockfd, int backlog);
```

marks the socket referred to by `sockfd` as a passive socket, i.e, a socket to be used to accept incoming connection requests using `accept()`

Return value: On success, zero is returned. On error, -1 is returned

Parameters:

`sockfd`: file descriptor that refers to a socket of type, e.g. `SOCK_STREAM`

`backlog`: defines the maximum length to which the queue of pending connections for `sockfd` may grow. If a connection request arrives when the queue is full, the client may receive an error with an indication of `ECONNREFUSED` or, if the underlying protocol supports retransmission, the request may be ignored so that a later reattempt succeeds.

¹If the backlog value is greater than the value in `/proc/sys/net/core/somaxconn`, then it is silently capped to that value. Since Linux 5.4, the default in this file is 4096; in earlier kernels, the default value is 128

The accept() API

```
#include <sys/socket.h>
int accept(int sockfd, struct sockaddr *addr,
           socklen_t *addrlen);
```

used with connection-oriented socket types (e.g. SOCK_STREAM). It extracts the first connection request on the queue of pending connections for the listening socket `sockfd`, creates a new connected socket, and returns a new file descriptor for it.

Return value: On success, returns a file descriptor for the accepted socket (a nonnegative integer). On error, -1 is returned

¹The newly created socket is not in the listening state. The original socket `sockfd` is unaffected

The `accept()` API

```
#include <sys/socket.h>
int accept(int sockfd, struct sockaddr *addr,
           socklen_t *addrlen);
```

Parameters:

sockfd: a socket created with `socket()`, bound to a local address with `bind()`, and is listening for connections using `listen()`

addr: a pointer to an address structure of the peer, actual structure depends on the socket address family

addrlen: a value-result argument; initialized to contain the size (in bytes) of the structure pointed to by **addr**; on return it will contain the actual size of the peer address

The `close()` API

```
#include <unistd.h>
int close(int fd);
```

closes a file descriptor, so that it no longer refers to any file and may be reused

Return value: returns zero on success. On error, -1 is returned

Parameters:

fd: closes the socket identified by the file descriptor

¹<https://man7.org/linux/man-pages/man2/close.2.html>

The read() API

```
#include <unistd.h>
```

```
ssize_t read(int fd, void *buf, size_t count);
```

attempts to read up to `count` bytes from file descriptor `fd` into the buffer starting at `buf`

Return value: On success, the number of bytes read is returned (zero indicates end of file), and the file position is advanced by this number. On error, -1 is returned

Parameters:

fd: a file descriptor to read from

buf: pointer to a buffer area (array) to read into

count: maximum number of bytes to read

¹<https://man7.org/linux/man-pages/man2/read.2.html>

²<https://man7.org/linux/man-pages/man2/recv.2.html>

The write() API

```
#include <unistd.h>
```

```
ssize_t write(int fd, void *buf, size_t count);
```

writes up to `count` bytes from buffer starting at `buf` into the file descriptor `fd`

Return value: On success, the number of bytes written is returned. On error, -1 is returned

Parameters:

`fd`: a file descriptor to write into

`buf`: pointer to a buffer area (array) to write from

`count`: maximum number of bytes to write

¹<https://man7.org/linux/man-pages/man2/write.2.html>

²<https://man7.org/linux/man-pages/man2/send.2.html>

Creating an Echo Server

`server1.c`

Testing the Echo Server

compile and run the server in a terminal; and leave it be

```
gcc server1.c -o server1 && ./server1
```

run a telnet¹ client in another terminal

```
telnet <server ip> <server port>
```

```
telnet localhost 54321
```

write anything in telnet

type “quit” (or send `ctrl+]`) to close the connection

¹Telnet is an application protocol used on the Internet or local area network to provide a bidirectional interactive text-oriented communication facility - WIKI